


Agentic AI


From open theme to working demo



What is "Agentic AI"?

An AI that **plans, acts, observes, and adapts** — not just a one-shot reply.

The Core Loop



Plan → Act → Observe → Adapt

Building Blocks

- **Tools** — search, calculators, APIs, code execution
- **Memory / state** — carrying info across steps or sessions
- **Orchestration** — single agent vs. multiple cooperating agents

Where Agentic AI Shows Up Today

- Coding agents
- Research / browsing agents
- Support automation
- Workflow tools



What Makes Agentic Systems Hard

- Reliability
- Evaluation
- Safety / guardrails
- Cost & latency



Two Kinds of Projects

1. **Building agentic AI** — the agent *is* the product
2. **Building a system that *uses* agentic AI** — it's one component in something bigger

The Complexity Ladder

- **Level 1** — One LLM call + one external tool
- **Level 2** — A loop that checks its own output and retries
- **Level 3** — Multiple tools, with decisions about which to use when
- **Level 4** — Multiple cooperating agents, or persistent memory over a long task

| *A well-executed Level 1 project beats a broken Level 4 one.*

Design Questions Worth Asking Early

- What's the agent's **goal** — how will it know it's done?
- What **tools/data** does it need?
- What happens when it **fails or gets stuck**?
- Where does this land on the **complexity ladder** — is that the right call?

Resources to Get Started

- Approved frameworks, APIs, docs
- API budgets / keys
- Compute limits



A Note on Scope

Narrow and well-defined beats ambitious and half-working.



**Now here's how you'll spend
the next 4 days →**

